

CS512 Assignment 4: Report  
Parthkumar Patel  
CWID: A20416508  
Semester: Fall 2018  
Department of Computer Science  
Illinois Institute of Technology  
November 03,2018

## ❖ Abstract

This programming assignment deals with implementation of a basic Convolutional Neural Network for classification. For, implementation it needs to use a GPU framework. Specifically, Keras or TensorFlow. After that we will train and evaluate a CNN to classify images of numbers in the MNIST dataset as either even or odd.

## ❖ Problem Statement

### Deliverables 1: Custom CNN

1. In this we need to train CNN using following configuration such as: use MNIST train set: 55,000 samples, add 2 Convolutional and Pooling layers in which Layer 1 should be 32 filters of kernel size 5x5 and Layer 2 should be 64 filters of kernel size 5x5, Pooling in both layers should down sample by a factor of 2, Dropout rate of 40%, use gradient descent optimization and a learning rate of 0.001, Train for 5 epochs
2. In this we need to report the training loss and accuracy for each step/iteration and plot the curve as our model trains. Also, we need to report the loss and accuracy value of the final training step
3. In this we need to report the loss, accuracy, precision and recall for MNIST test set

### Deliverables 2: Parameter Tuning

Evaluate different variations of the basic network as described in the question and measure performance

### Deliverable 3: Application

In this we need write a program using our pretrained custom CNN which should do the following:

1. Accept as input an image of a handwritten digit
2. Using OpenCV do some basic image pre-processing to prepare the image for our CNN
3. Use our CNN to classify the binary image
4. Our program should continuously request the path to an image file, process the image, and output to the console that class (even/odd) of the image
5. Program should terminate when 'q' or ESC key is entered

## ❖ Proposed Solutions

First, we will train and evaluate a CNN to classify images of numbers in the MNIST dataset as either even or odd. Images in the MNIST dataset are labelled by the corresponding integer. For example, a handwritten image of the digit '7' is labelled as the integer 7. Hence, you will have to map the integer labelling (10 classes) of the images to a binary labelling (even/odd).

After, that we will build the linear stack of layers using sequential model where we will add 2 layers where first one is 32 filter of kernel size  $5 \times 5$  and second one 64 filters of kernel size  $5 \times 5$ . Now, we will down sample it by factor of 2 and set the other parameters and will use gradient decent optimization and a learning rate of 0.001 and train it for 5 epochs.

After training it for 5 epochs, now we will train it for final training step and calculate the loss, accuracy, precision and recall. Now, we will test it using 10000 images and we will get the result of the above listed parameter for MNSIT test images. Now , we will save that created model using keras's `save.model` function and load it into another file to test our handwritten images to check whether it will classify it into odd or even image.

## ❖ Results

### Deliverables– 1

#### 1. Trained CNN using the configuration mentioned in the question

```
#Training for 5 epochs
epochs = 5

# building a linear stack of layers with the sequential model
model = Sequential()
model.add(Conv2D(32, kernel_size=(5,5),input_shape=(28, 28, 1), activation = 'relu')) #Layer 1 consist of 32 filters of kernel s
model.add(MaxPooling2D(pool_size=(2, 2))) #pooling and downsample by a factor of 2
model.add(Conv2D(64, kernel_size=(5,5), activation = 'relu')) #Layer 2 consist of 64 filters of kernel size 5 X 5
model.add(MaxPooling2D(pool_size=(2, 2))) #pooling and downsample by a factor of 2
model.add(Dropout(0.4)) #droupout rate of 40%

# Flatten Layer
model.add(Flatten())

# Use final dense layer to get two outputs
model.add(Dense(1))

#using sigmoid activation
model.add(Activation("sigmoid"))

#using gradient descent optimization and a learning rate of 0.001
sgd = optimizers.SGD(lr=0.001)

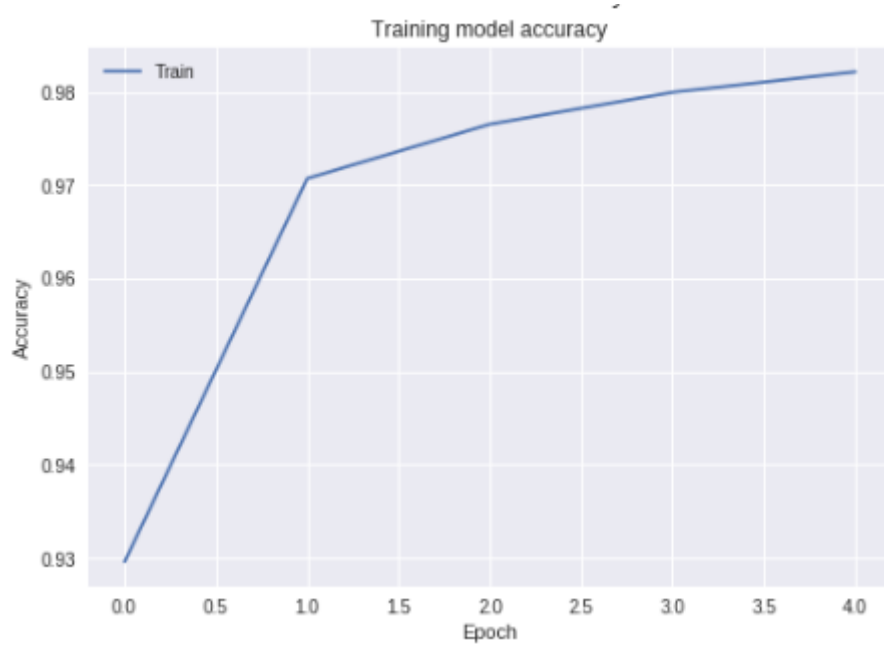
# compiling the sequential model
model.compile(loss='binary_crossentropy', optimizer=sgd, metrics=['accuracy', precision, recall])

# training the model and saving metrics in history
history = model.fit(x_train, y_train, epochs= epochs , validation_data=(x_test, y_test))
```

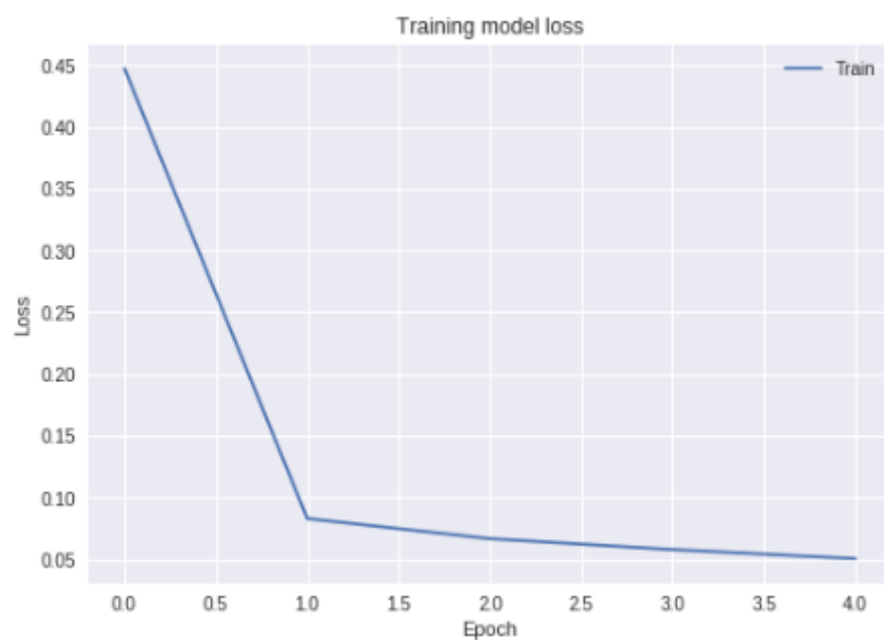
#### 2.

##### a. training loss and accuracy for each step/iteration and plot

```
Train on 55000 samples, validate on 10000 samples
Epoch 1/5
55000/55000 [=====] - 15s 264us/step - loss: 0.3316 - acc: 0.9301
Epoch 2/5
55000/55000 [=====] - 14s 257us/step - loss: 0.0879 - acc: 0.9679
Epoch 3/5
55000/55000 [=====] - 14s 258us/step - loss: 0.0657 - acc: 0.9767
Epoch 4/5
55000/55000 [=====] - 14s 258us/step - loss: 0.0581 - acc: 0.9799
Epoch 5/5
55000/55000 [=====] - 14s 259us/step - loss: 0.0503 - acc: 0.9827
```



Plot of training model accuracy



Plot of training model loss

b. Loss and accuracy value of the final training step

55000/55000 [=====] - 6s 118us/step

Loss and accuracy of final training step

Train loss: 0.028436766319383276

Train accuracy: 0.9904

3.

a. Report the loss, accuracy, precision and recall for MNIST test set

Evaluation of Testing set:

10000/10000 [=====] - 1s 120us/step

Test Loss:8.178321719360351 Test Accuracy:0.4926 Test Precision:0.8700974410057067 Test Recall:0.025089589014649392

b. For each epoch trained compute above mentioned metrics

Train on 55000 samples, validate on 10000 samples

Epoch 1/5

55000/55000 [=====] - 15s 271us/step - loss: 8.1936 - acc: 0.4915 - precision: 0.7850 - recall: 0.0453 - val\_loss: 8.1783 - val\_acc: 0.4926 - val\_precision: 0.8383 - val\_recall: 0.0330

Epoch 2/5

55000/55000 [=====] - 14s 263us/step - loss: 8.1953 - acc: 0.4915 - precision: 0.8454 - recall: 0.0308 - val\_loss: 8.1783 - val\_acc: 0.4926 - val\_precision: 0.8531 - val\_recall: 0.0308

Epoch 3/5

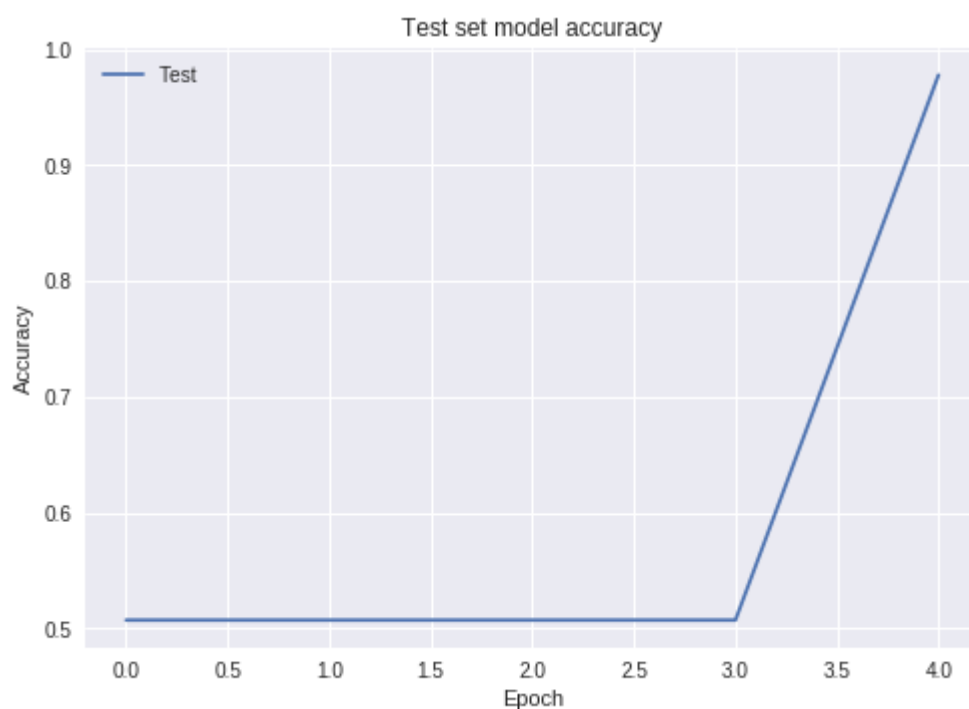
55000/55000 [=====] - 15s 264us/step - loss: 8.1953 - acc: 0.4915 - precision: 0.8567 - recall: 0.0299 - val\_loss: 8.1783 - val\_acc: 0.4926 - val\_precision: 0.8628 - val\_recall: 0.0300

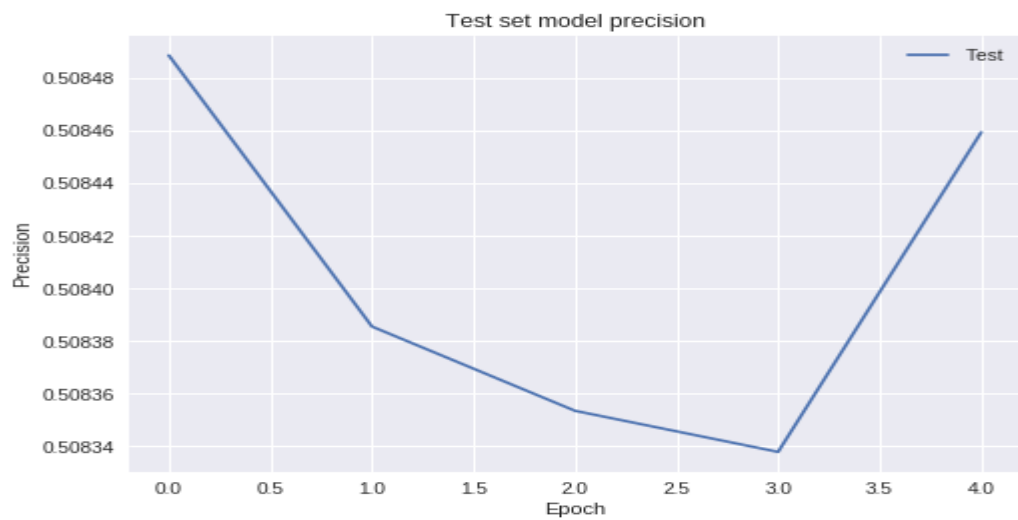

Epoch 4/5

55000/55000 [=====] - 15s 266us/step - loss: 8.1953 - acc: 0.4915 - precision: 0.8644 - recall: 0.0296 - val\_loss: 8.1783 - val\_acc: 0.4926 - val\_precision: 0.8663 - val\_recall: 0.0294

Epoch 5/5

55000/55000 [=====] - 15s 268us/step - loss: 8.1953 - acc: 0.4915 - precision: 0.8663 - recall: 0.0294 - val\_loss: 8.1783 - val\_acc: 0.4926 - val\_precision: 0.8676 - val\_recall: 0.0294





## Deliverables 2:

### ➤ Learning Rate: 0.01

Train on 55000 samples, validate on 10000 samples

Epoch 1/5

55000/55000 [=====] - 15s 268us/step - loss: 7.8351 - acc: 0.5085 - precision: 0.5082 - recall: 0.9994

Epoch 2/5

55000/55000 [=====] - 14s 257us/step - loss: 7.8364 - acc: 0.5085 - precision: 0.5079 - recall: 1.0000

Epoch 3/5

55000/55000 [=====] - 14s 258us/step - loss: 7.8364 - acc: 0.5085 - precision: 0.5080 - recall: 1.0000

Epoch 4/5

55000/55000 [=====] - 14s 258us/step - loss: 7.8364 - acc: 0.5085 - precision: 0.5084 - recall: 1.0000

Epoch 5/5

55000/55000 [=====] - 14s 258us/step - loss: 7.8364 - acc: 0.5085 - precision: 0.5085 - recall: 1.0000

55000/55000 [=====] - 6s 116us/step

Loss and accuracy of final training step

Train loss: 7.83640719687722

Train accuracy: 0.5084545454502106

Evaluation of Testing set:

10000/10000 [=====] - 1s 116us/step

Test Loss:7.85321912612915 Test Accuracy:0.5074 Test Precision:0.5083225183486938 Test Recall:1.0

When we changed the learning rate of the model from 0.001 to 0.01 then the loss is increased and accuracy of the model is decreased from 0.98 to 0.50.

### ➤ Dropout – 30%

Train on 55000 samples, validate on 10000 samples

Epoch 1/5

55000/55000 [=====] - 15s 267us/step - loss: 6.1876 - acc: 0.5991 - precision: 0.7635 - recall: 0.2862

Epoch 2/5

55000/55000 [=====] - 14s 257us/step - loss: 0.1107 - acc: 0.9604 - precision: 0.5450 - recall: 0.6690

Epoch 3/5

55000/55000 [=====] - 14s 257us/step - loss: 0.0730 - acc: 0.9745 - precision: 0.5248 - recall: 0.8096

Epoch 4/5

55000/55000 [=====] - 14s 256us/step - loss: 0.0571 - acc: 0.9804 - precision: 0.5191 - recall: 0.8659

Epoch 5/5

55000/55000 [=====] - 14s 256us/step - loss: 0.0489 - acc: 0.9836 - precision: 0.5163 - recall: 0.8965

55000/55000 [=====] - 6s 116us/step

Loss and accuracy of final training step

Train loss: 0.027314712133309382

Train accuracy: 0.990909090909091

Evaluation of Testing set:

10000/10000 [=====] - 1s 118us/step

Test Loss:0.0282995004942175 Test Accuracy:0.9894 Test Precision:0.5139981072425842 Test Recall:0.9228913857460022

Training model accuracy:

When we changed the drop out from 40% to 30% then the precision is decreased from around 0.87 to 0.51.

### ➤ Epochs: 10

```
Epoch 4/10
55000/55000 [=====] - 14s 257us/step - loss: 0.0651 - acc: 0.9775 - precision: 0.5083 - recall: 1.0000
Epoch 5/10
55000/55000 [=====] - 14s 255us/step - loss: 0.0568 - acc: 0.9804 - precision: 0.5085 - recall: 1.0000
Epoch 6/10
55000/55000 [=====] - 14s 255us/step - loss: 0.0511 - acc: 0.9822 - precision: 0.5085 - recall: 1.0000
Epoch 7/10
55000/55000 [=====] - 14s 256us/step - loss: 0.0471 - acc: 0.9831 - precision: 0.5085 - recall: 1.0000
Epoch 8/10
55000/55000 [=====] - 14s 255us/step - loss: 0.0428 - acc: 0.9856 - precision: 0.5084 - recall: 1.0000
Epoch 9/10
55000/55000 [=====] - 14s 256us/step - loss: 0.0418 - acc: 0.9855 - precision: 0.5084 - recall: 1.0000
Epoch 10/10
55000/55000 [=====] - 14s 256us/step - loss: 0.0386 - acc: 0.9864 - precision: 0.5084 - recall: 1.0000
55000/55000 [=====] - 6s 115us/step
```

Loss and accuracy of final training step  
Train loss: 0.020452919308942826  
Train accuracy: 0.9938727272727272

Evaluation of Testing set:

```
10000/10000 [=====] - 1s 115us/step
Test Loss:0.02530967883798294 Test Accuracy:0.9919 Test Precision:0.5083759575843811 Test Recall:0.9999916866302491
```

### ➤ Epochs: 3

```
Train on 55000 samples, validate on 10000 samples
Epoch 1/3
55000/55000 [=====] - 15s 268us/step - loss: 0.5155 - acc: 0.9223 - precision: 0.5277 - recall: 0.9905
Epoch 2/3
55000/55000 [=====] - 14s 257us/step - loss: 0.0906 - acc: 0.9681 - precision: 0.5113 - recall: 0.9984
Epoch 3/3
55000/55000 [=====] - 14s 256us/step - loss: 0.0706 - acc: 0.9751 - precision: 0.5100 - recall: 0.9991
55000/55000 [=====] - 6s 116us/step
```

Loss and accuracy of final training step  
Train loss: 0.036894038953395054  
Train accuracy: 0.9875636363636363

Evaluation of Testing set:

```
10000/10000 [=====] - 1s 117us/step
Test Loss:0.0382737568805227 Test Accuracy:0.986 Test Precision:0.509482096195221 Test Recall:0.9994521673202514
```



➤ Optimizer = Adam

```
📁 Train on 55000 samples, validate on 10000 samples
Epoch 1/5
55000/55000 [=====] - 17s 304us/step - loss: 0.3618 - acc: 0.9476 - precision: 0.5319 - recall: 0.9932
Epoch 2/5
55000/55000 [=====] - 16s 288us/step - loss: 0.0647 - acc: 0.9789 - precision: 0.5119 - recall: 0.9989
Epoch 3/5
55000/55000 [=====] - 16s 284us/step - loss: 0.0599 - acc: 0.9803 - precision: 0.5106 - recall: 0.9994
Epoch 4/5
55000/55000 [=====] - 16s 284us/step - loss: 0.0557 - acc: 0.9823 - precision: 0.5101 - recall: 0.9996
Epoch 5/5
55000/55000 [=====] - 16s 284us/step - loss: 0.0541 - acc: 0.9830 - precision: 0.5096 - recall: 0.9997
55000/55000 [=====] - 6s 115us/step

Loss and accuracy of final training step
Train loss: 0.028247765320996668
Train accuracy: 0.9906909090909091

Evaluation of Testing set:


10000/10000 [=====] - 1s 116us/step
Test Loss:0.027219966612639836 Test Accuracy:0.9905 Test Precision:0.509289197921753 Test Recall:0.9997444877624512
```

**Deliverable 3:**

```
Using TensorFlow backend.
enter name of the image with .jpg extension or enter 'q' for exit:
6.jpg
[[0.15105172]]
even
>>>
```



```
Using TensorFlow backend.  
enter name of the image with .jpg extension or enter 'q' for exit:  
5.jpg  
[[0.9954639]]  
odd  
>>>
```



5

```
Using TensorFlow backend.  
enter name of the image with .jpg extension or enter 'q' for exit:  
3.jpg  
[[0.99947745]]  
odd  
>>>
```



3

```
Using TensorFlow backend.  
enter name of the image with .jpg extension or enter 'q' for exit:  
7.jpg  
[[0.98492056]]  
odd  
>>>
```

7